


Statistics Without Borders

# Deep Learning II

Luca Alberto Rizzo




# Course outline

1. Applications of Deep Learning
2. What to do before you start your DL project
3. Introduction to torchvision
4. Opening (a bit) the DL blackbox
5. Additional reading
6. Presentation of the demonstration




# Applications of Deep Learning



# Applications of Deep Learning: exponential growth

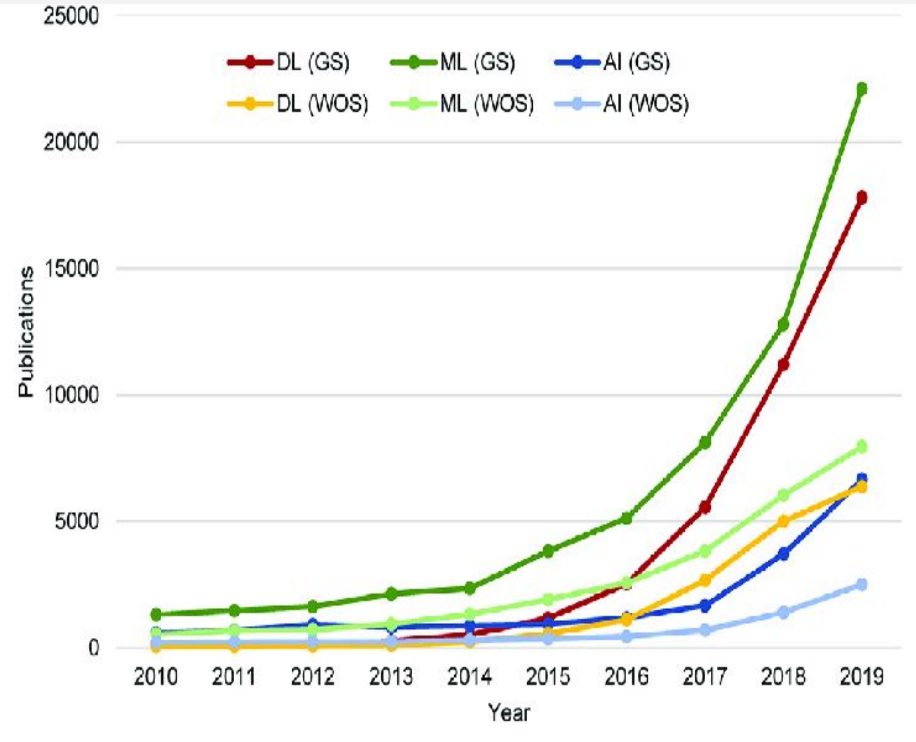
PRECEDENCE  
RESEARCH

DEEP LEARNING MARKET SIZE 2023 TO 2033 (USD BILLION)



Source: <https://www.precedenceresearch.com/deep-learning-market>

[Precedence research](https://www.precedenceresearch.com)



[A Bird's-Eye View of Deep Learning in Bioimage Analysis. E. Meijering](#)



STATISTICS WITHOUT  
BORDERS

@SWBprobono

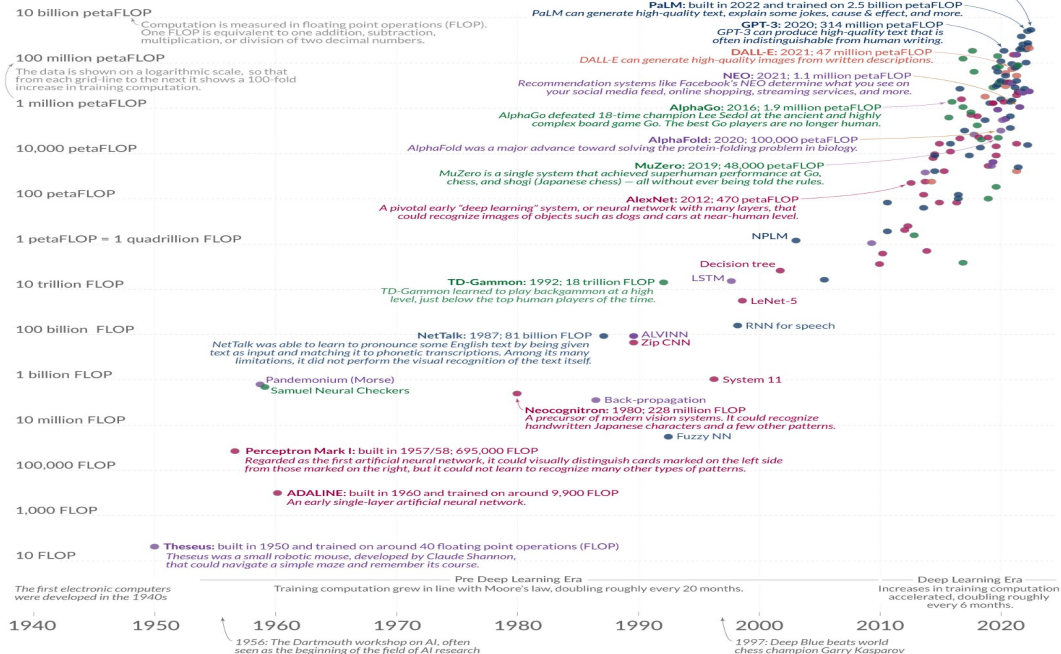
[StatisticsWithoutBorders.org](https://www.StatisticsWithoutBorders.org)

# Applications of Deep Learning: why now?

**The rise of artificial intelligence over the last 8 decades: As training computation has increased, AI systems have become more powerful**

The color indicates the domain of the AI system: ● Vision ● Games ● Drawing ● Language ● Other

Shown on the vertical axis is the **training computation** that was used to train the AI systems.



The data on training computation is taken from Sevilla et al. (2022) - Parameter, Compute, and Data Trends in Machine Learning. It is estimated by the authors and comes with some uncertainty. The authors expect the estimates to be correct within a factor of two.

OurWorldInData.org - Research and data to make progress against the world's largest problems. Licensed under CC-BY by the authors Charlie Giattino, Edouard Mathieu, and Max Roser

Our World in Data

Simultaneously exponential growth of:

- computational power
- data size

Deep Learning performs **better** with a lot of data but requires a lot of computational power

[Source : Our World in Data](https://ourworldindata.org)



STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org

# Applications of Deep Learning: examples

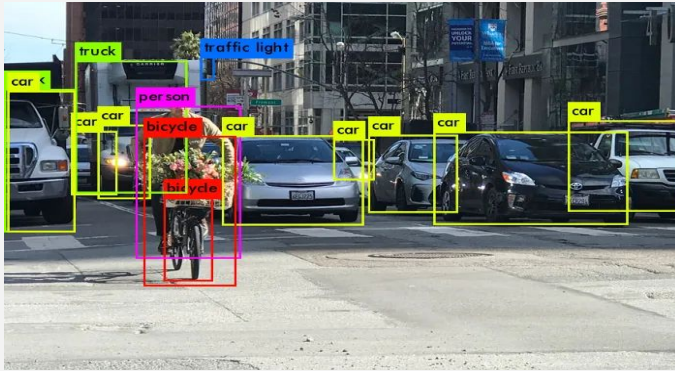
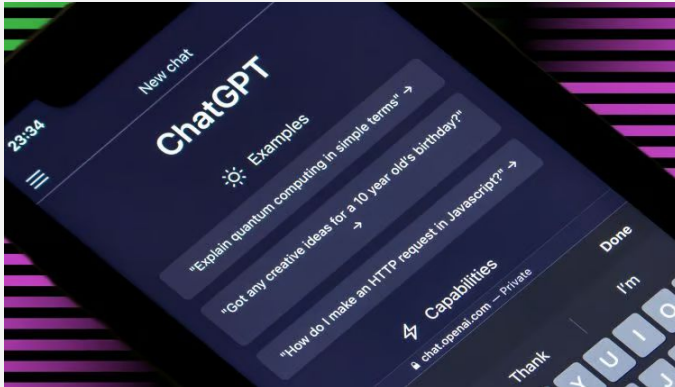
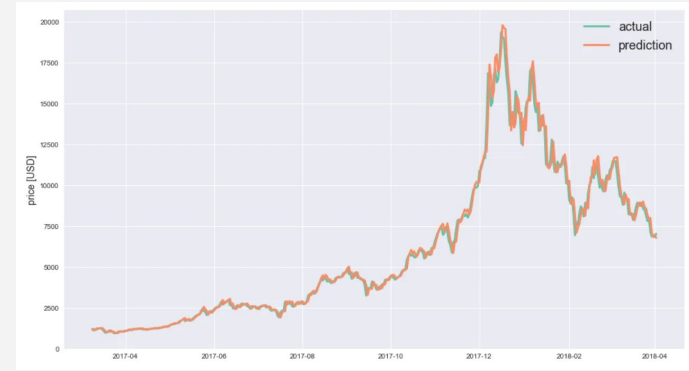


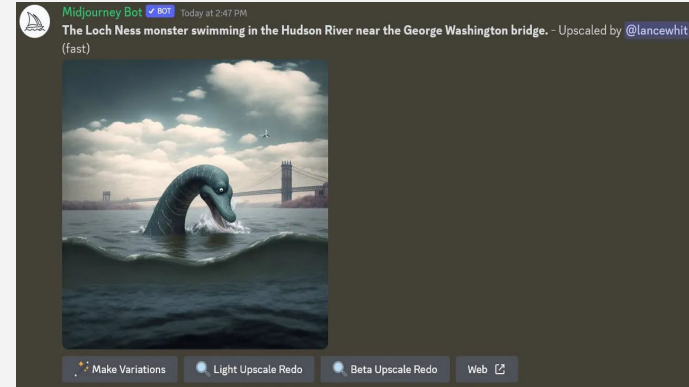
Image recognition with YOLO



Natural language processing (e.g. ChatGPT)




Portfolio prediction in Finance



Natural language processing (e.g. ChatGPT)



# What to do before you start your DL project







# What to do before you start: how complex is your problem?

[Review of tool condition monitoring in machining and opportunities for deep learning](#)

[Review of tool condition monitoring in machining and opportunities for deep learning](#)



If you have enough unstructured data, it might be a good idea to use Deep Learning



STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org



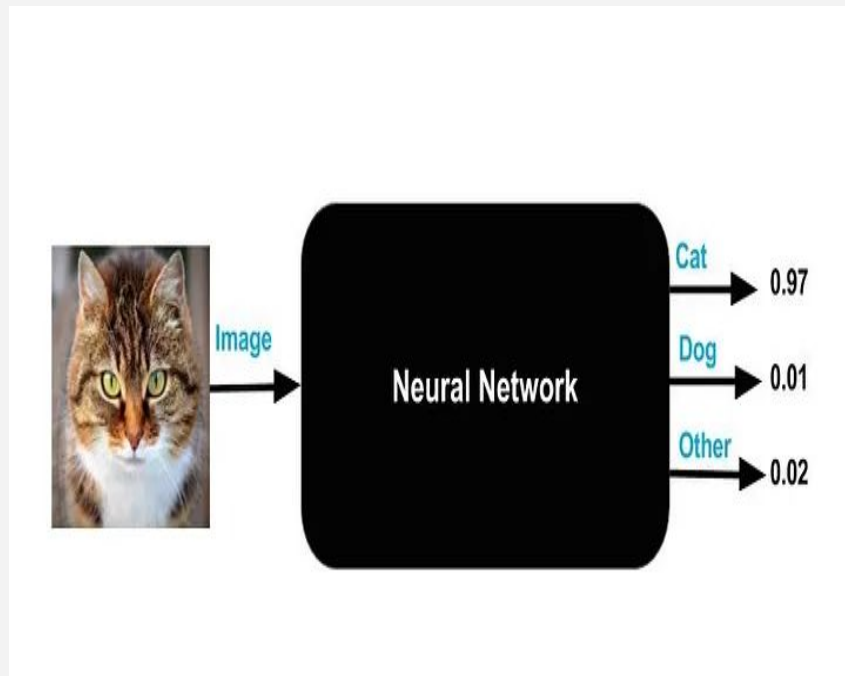
# What to do before you start: when you might not use DL

Deep Learning is:

- Expensive
- Not interpretable (even more than ML)
- Very hard to optimize
- Not so great with tabular data

However:

- It works extremely well for unstructured data!
- Plenty of resources around



**PLEASE RESIST THE HYPE!**

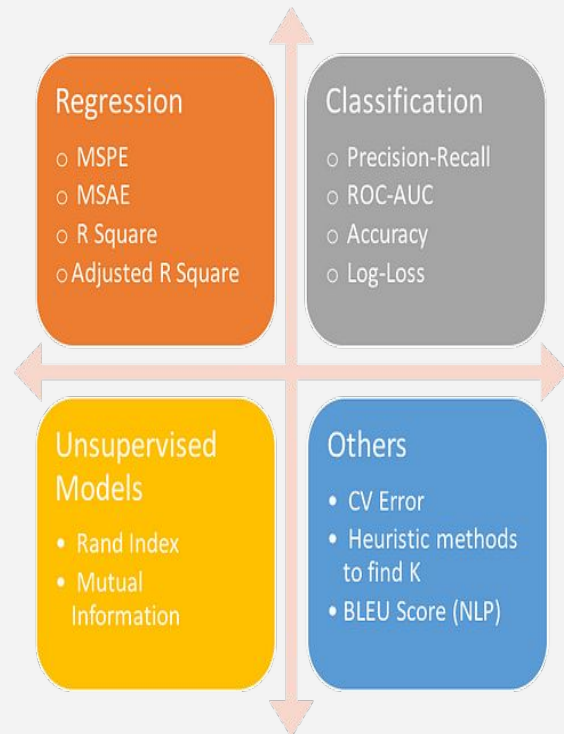
[How to explore Neural networks, the black box ?](#)



# What to do before you start: how to choose the right metric

## Your model is as good as your metric

- Talk with the final user (several times)
- Think about worst and best case scenarios
- Connect final metric with training metrics



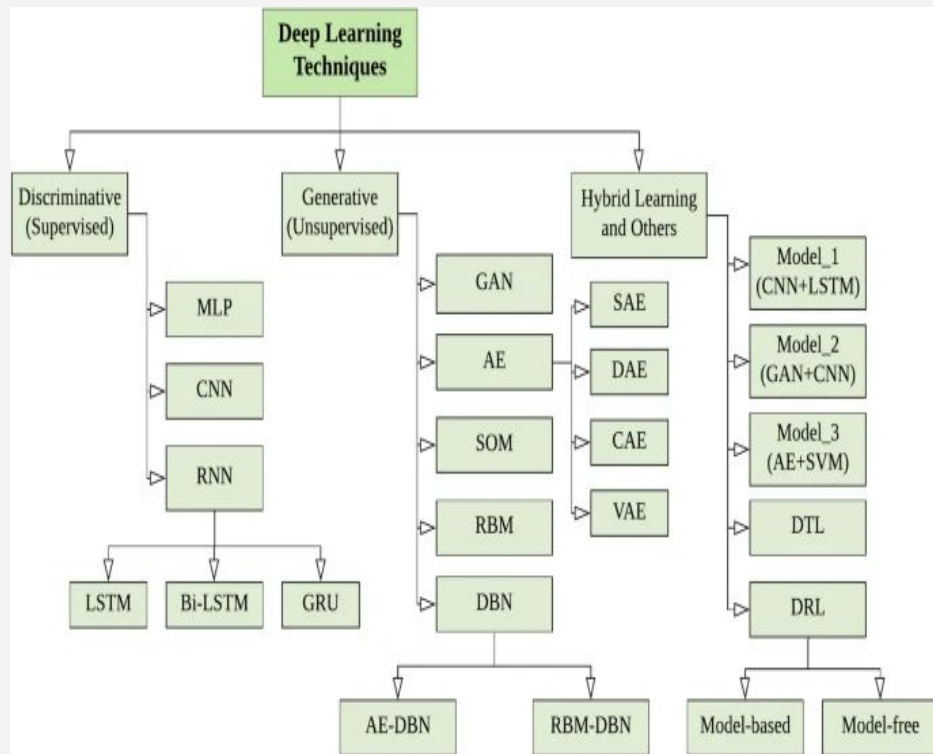
[Source: Performance Metrics for Classification Problems \(Kaggle\)](#)



# What to do before you start: how to choose the right model

Most models are easy to implement

- **Ask around**
- **Read literature / blogs** (e.g. [machine learning mastery](#))
- **Try on smaller version** of your dataset



[Deep Learning: A Comprehensive Overview](#)



STATISTICS WITHOUT  
BORDERS

@SWBprobono

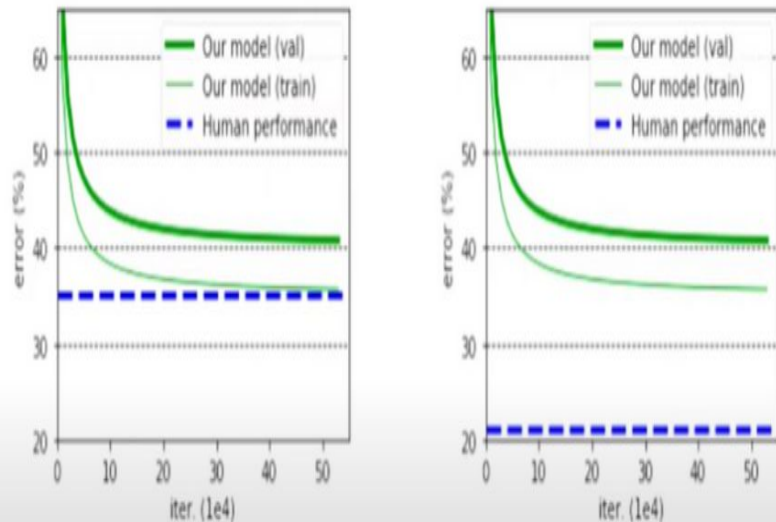
StatisticsWithoutBorders.org

# What to do before you start: how to choose the right baseline

## Baselines are extremely important

- **Better than random** (e.g. 50% accuracy for binary classification)
- **DL must be better** than simpler (ML) models
- **DL could be better** than humans

## Same model, different baseline → different next steps




[Source: how to set a baseline \(The Full Stack\)](#)



# What to do before you start: transfer learning

Larger, already trained models can be used for your problem

Plenty of models are easily available on Pytorch



[Source: Geeks for Geeks](#)

## Classification

The following classification models are available, with or without pre-trained weights:

- AlexNet
- ConvNeXt
- DenseNet
- EfficientNet
- EfficientNetV2
- GoogleNet
- Inception V3
- MaxViT
- MNASNet
- MobileNet V2
- MobileNet V3
- RegNet
- ResNet
- ResNeXt
- ShuffleNet V2
- SqueezeNet
- SwinTransformer
- VGG
- VisionTransformer
- Wide ResNet

## Object Detection

The following object detection models are available, with or without pre-trained weights:

- Faster R-CNN
- FCOS
- RetinaNet
- SSD
- SSDlite

## Semantic Segmentation

### WARNING

The segmentation module is in Beta stage, and backward compatibility is not guaranteed.

The following semantic segmentation models are available, with or without pre-trained weights:

- DeepLabV3
- FCN
- LRASPP

## Video Classification

### WARNING

The video module is in Beta stage, and backward compatibility is not guaranteed.

The following video classification models are available, with or without pre-trained weights:

- Video SwiT
- Video ResNet
- Video SSD
- Video SwinTransformer

**TRY TO NOT REINVENT THE WHEEL!**




STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org

# Introduction to torchvision



# Torchvision: why torchvision?

## **torchvision**

This library is part of the **PyTorch** project. PyTorch is an open source machine learning framework.

- Easy to install and use
- Plenty of models available
- Plenty of tools (as in Pytorch)
- Seamless integration with video and audio tools



## ECOSYSTEM TOOLS

Tap into a rich ecosystem of tools, libraries, and more to support, accelerate, and explore AI development.

Join the Ecosystem

[Pytorch ecosystem](https://pytorch.org/ecosystem/)



STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org



# Torchvision: MNIST

MNIST is a large handwritten digits database, commonly used for computer vision

- Published by the National Institute of Standards and Technology in 1994
- 6000 training and 10000 images
- 128x128 images grayscale to 28x28

**It has been used widely to test image recognition algorithm**

Torchvision implements MNIST natively, along with several other important datasets



[Source: wikipedia page](#)

## Built-in datasets

All datasets are subclasses of `torch.utils.data.Dataset` i.e. they have `__getitem__` and `__len__` methods implemented. Hence, they can all be passed to a `torch.utils.data.DataLoader` which can load multiple samples in parallel using `torch multiprocessing` workers. For example:

```
imagenet_data = torchvision.datasets.ImageNet('path/to/imagenet_root/')
data_loader = torch.utils.data.DataLoader(imagenet_data,
    batch_size=4,
    shuffle=True,
    num_workers=args.nThreads)
```

All the datasets have almost similar API. They all have two common arguments: `transform` and `target_transform` to transform the input and target respectively. You can also create your own datasets using the provided [base classes](#).

## Image classification

```
Caltech101(root[, target_type, transform, ...])
```

Caltech 101 Dataset.

```
Caltech256(root[, transform, ...])
```

Caltech 256 Dataset.



# Torchvision: load MNIST

With one-liner we can:

- Define image transformation
- Download (or load locally) MNIST
- Prepare dataset for training

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

batch_size = 10

trainset = torchvision.datasets.MNIST(root='./data', train=True,
                                       download=True, transform=transform)

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.MNIST(root='./data', train=False,
                                       download=True, transform=transform)

testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)
```

Dataset stores the samples and their corresponding labels,  
and DataLoader wraps an iterable around the Dataset to enable easy access to the samples.



# Torchvision: why torchvision?

Let us define a simple neural network:

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output
```

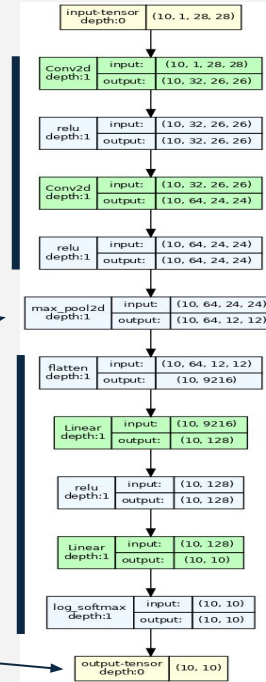
And visualize with torchview:

2 Convolutions

Pooling (reduce dimensionality)

Flatten data

Compute class probabilities



Moder neural networks architectures can be complicated (not necessarily better)



# Torchvision: criterion and optimizer

## Choose training criterion and optimizer

How do we choose the best model?

```
criterion = nn.CrossEntropyLoss()
```

### Loss Functions

<code>nn.L1Loss</code>	Creates a criterion that measures the mean absolute error (MAE) between each element in the input $x$ and target $y$ .
<code>nn.MSELoss</code>	Creates a criterion that measures the mean squared error (squared L2 norm) between each element in the input $x$ and target $y$ .
<code>nn.CrossEntropyLoss</code>	This criterion computes the cross entropy loss between input logits and target.
<code>nn.CTCLoss</code>	The Connectionist Temporal Classification loss.
<code>nn.NLLLoss</code>	The negative log likelihood loss.
<code>nn.PoissonNLLLoss</code>	Negative log likelihood loss with Poisson distribution of target.
<code>nn.GaussianNLLLoss</code>	Gaussian negative log likelihood loss.

How do we find the best model?

```
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

### Algorithms

<code>Adadelta</code>	Implements Adadelta algorithm.
<code>Adagrad</code>	Implements Adagrad algorithm.
<code>Adam</code>	Implements Adam algorithm.
<code>AdamW</code>	Implements AdamW algorithm.
<code>SparseAdam</code>	SparseAdam implements a masked version of the Adam algorithm suitable for sparse gradients.
<code>Adamax</code>	Implements Adamax algorithm (a variant of Adam based on infinity norm).
<code>ASGD</code>	Implements Averaged Stochastic Gradient Descent.

A large family of both exist: results might depend drastically on criterion (less on optimizer)



# Torchvision: train the model

```
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.3f}')
            running_loss = 0.0

print('Finished Training')
```

```
[1, 2000] loss: 0.208
[1, 4000] loss: 0.101
[1, 6000] loss: 0.072
[2, 2000] loss: 0.055
[2, 4000] loss: 0.053
[2, 6000] loss: 0.045
Finished Training
```

Looping several times on the dataset allows the fit to converge better

Pytorch accepts data as trainloader to optimizer training

Setting optimizer parameter to zero

Compute outputs with **current** model

Compute loss with **current** model

Pytorch's magic to tell the optimizer how much and in which direction it should go

Printing training loss every 2000 iterations



# Torchvision: testing model

Compute results on test statistics and use them for the confusion matrix

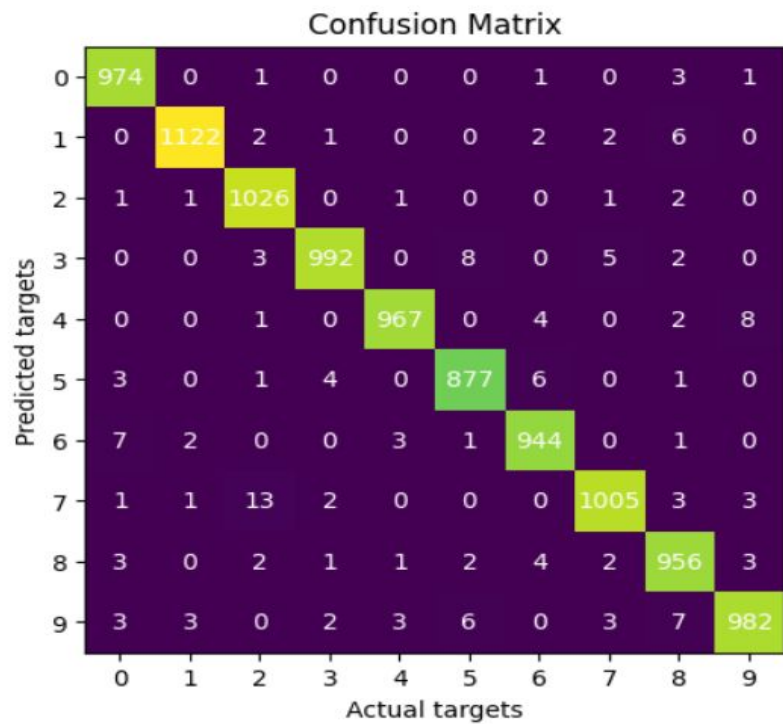
```
# Test function
def testing_accuracy(model, data_loader):
    model.eval()
    test_loss = 0
    device = 'cpu'

    y_pred = []
    y_actu = []
    with torch.no_grad():
        for data, target in data_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
            pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
            y_pred.extend(torch.flatten(pred).tolist())
            y_actu.extend(target.tolist())

    y_pred = pd.Series(y_pred, name='Actual')
    y_actu = pd.Series(y_actu, name='Predicted')
    cm = pd.crosstab(y_actu, y_pred)
    correct = sum([cm.iloc[i,i] for i in range(len(cm))])

    test_loss /= len(data_loader.dataset)
    accuracy = 100*correct/len(data_loader.dataset)

    return(test_loss, accuracy, cm)
```





# Torchvision: save and load models

Our model can be saved and loaded easily

```
PATH = './my_minst_model_net.pth'  
torch.save(net.state_dict(), PATH)
```

```
model = torch.load(PATH)
```

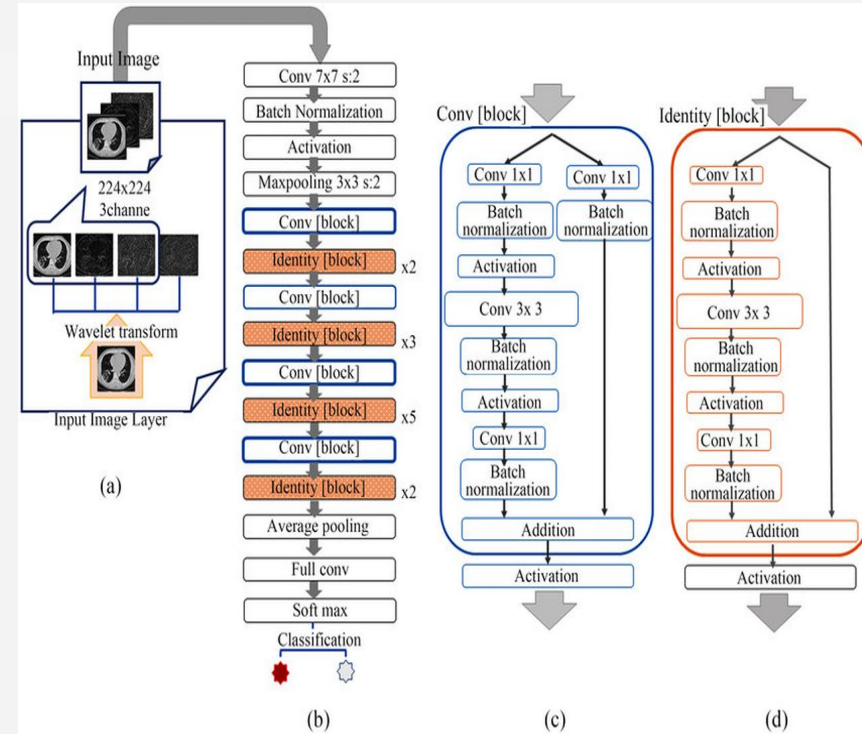
Large models can be loaded and modified too

```
resnet = models.resnet50(pretrained=True)
```

```
def change_layers(model):  
    model.conv1 = nn.Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
    model.fc = nn.Linear(2048, 10, bias=True)  
    return model
```

```
change_layers(resnet)
```

Out of the box accuracy is 8% ...



[Original architecture of ResNet50. Source: Wisdom ML](#)





# Torchvision: save and load models

Fine tuning is needed to achieve better performances

#FINE TUNING

```
optimizer_resnet = optim.SGD(resnet.parameters(), lr=0.001, momentum=0.9)

for epoch in range(2): # loop over the dataset multiple times

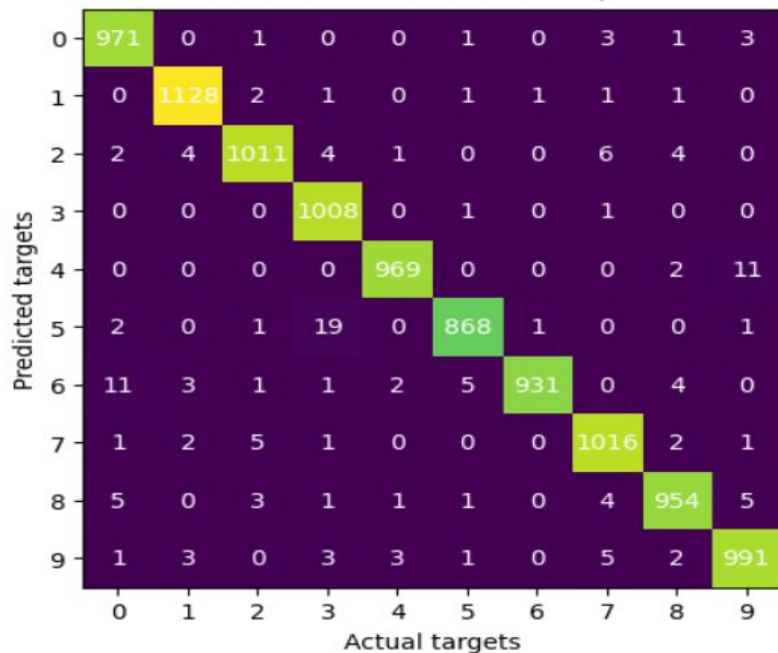
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer_resnet.step()

    # print statistics
    running_loss += loss.item()
    if i % 2000 == 1999: # print every 2000 mini-batches
        print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.3f}')
        running_loss = 0.0
```


Confusion Matrix for custom model (ACC = 99.99)



Slightly better acc. but model almost 25 bigger



# Opening (a bit) the blackbox





# Opening the black box: what is the network really learning

Saliency methods can be used to understand better deep learning results:

1. Perform forward pass
2. Compute gradient of class score with respect input pixels

$$E_{grad}(I_0) = \frac{\delta S_c}{\delta I} \Big|_{I=I_0}$$

3. Visualize the gradients



DL is learning differently than us!

[Saliency vs Solob methods \(source Xplique documentation\)](#)



STATISTICS WITHOUT  
BORDERS


@SWBprobono

StatisticsWithoutBorders.org


# Opening the black box: why so many parameters

Billions of parameter for LLMs

Parameters of transformer-based language models



Classical vs DL risk curve



**DL is overfit but acts differently than standard ML models: nobody knows why...**

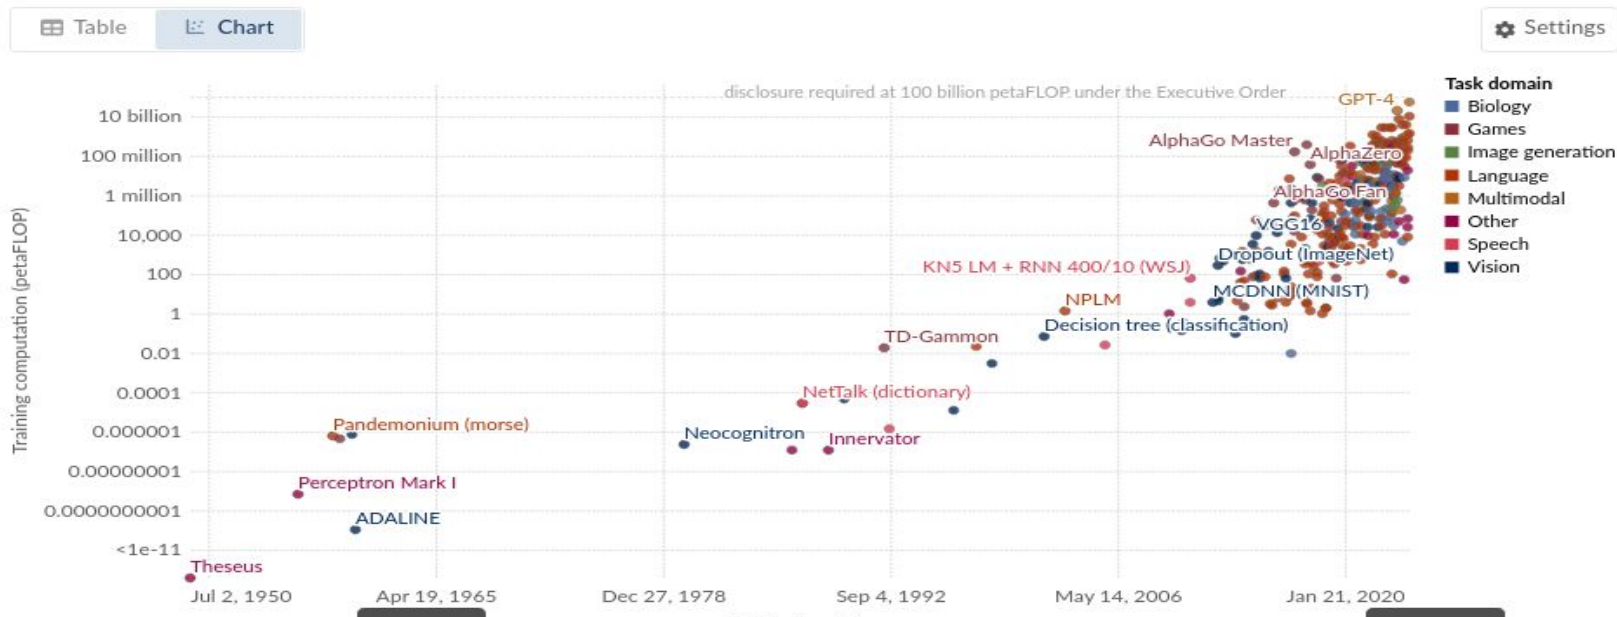


# Opening the black box: why so many parameters (energy)

## Computation used to train notable artificial intelligence systems

Computation is measured in total petaFLOP, which is  $10^{15}$  floating-point operations. Estimated from AI literature, albeit with some uncertainty. Estimates are expected to be accurate within a factor of 2, or a factor of 5 for recent undisclosed models like GPT-4.

Our World in Data



DL is particularly computationally expensive and becoming more so



STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org

# Opening the black box: ML vs DL


	Machine Learning	Deep Learning
Data Volume	Fewer, structured data	More, structured data
Computational Cost	Faster, lighter model	Very expensive training
Feature engineering	Explicitly done by humans	Often DL handles automatically
Interpretability	Usually easier	Extremely hard to interpret
Applications examples	Customer segmentation, Recommender systems, Fraud detection, ...	Computer Vision, Natural Language Processing, ...

**The differences are blurry and DL can be used for all applications**

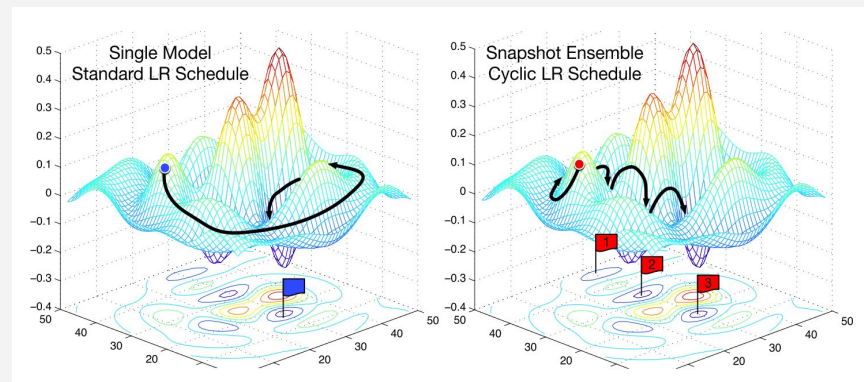





# Opening the black box: topics we could not possibly cover



Reinforcement learning



Role of optimizers




Generative AI (deep learning)





# Conclusions and further reading







# Conclusions

- Deep learning is a powerful and “easy-to-use” tool: please don’t reinvent the wheel!
- Deep learning should be used with care: complex, unstructured data work better
- Deep learning usage exploded but there are still a lot of unknowns
- Pytorch (and its sister libraries) is a powerful tool
- Deep learning is still a black box: use at your own risk!
- Reinforcement learning, Generative AI, ...



# Further reading: free resources




STATISTICS WITHOUT  
BORDERS

@SWBprobono

StatisticsWithoutBorders.org

# Presentation of the demonstration




# Presentation of the demonstration: up to you now!

- Install torchvision on your PC
- Load and explore MNIST dataset
- Built your first CNN
- Train your first CNN
- Explore your CNN results
- Experiments with different architectures
- Load pretrained network
- Fine tune a pretrained network and compare
- Is my network overtrained? Train-Test loss
- ...
- **WORK IN PROGRESS**



# Don't hesitate to contact me!



**deep-learning-II-lagos-lectures**



**Luca Alberto Rizzo**



**luca.alb.rizzo@gmail.com**




**STATISTICS WITHOUT  
BORDERS**

**@SWBprobono**

**StatisticsWithoutBorders.org**

Thank you for your attention






# SWB is grateful for the partnership of University of Lagos Department of Statistics

<https://StatisticsWithoutBorders.org>



[StatisticsWithoutBorders@gmail.com](mailto:StatisticsWithoutBorders@gmail.com)



**STATISTICS WITHOUT BORDERS**